

Latency vs budget per K — higher rate diverges sooner as K grows (phantom load)

